
ECSE 551 A2: Bernoulli Naïve Bayes Classification

Ashley Meagher (260822930), Charles Sirois (261158513)

Abstract

This project classifies posts on Reddit into four subreddits, London, Montreal, Paris, and Toronto, using the Bernoulli Naïve Bayes classification. The project requires text classification to take the text from the posts as input and assign it to a label (the subreddit). In the project, the first dataset contains posts for training and the second dataset contains posts to test. To improve the accuracy of the text classification, the features are preprocessed using stop-words, lemmatization, n-grams, language detection, and feature selection. Scikit-learn Decision Tree and SVM algorithms are used to compare the classification method, specifically using k-fold cross-validation. After feature manipulation and through cross-validation, it was found that the SVM classifier had the best results of all classifiers, with an accuracy of 72.2% and performance accuracy of 71.1% in the Kaggle competition.

1 Introduction

The goal of this project is to gain experience working with text classification in machine learning. Text classification is an important tool that filters text data, used in many real-life applications such as phishing emails and Google searches. This project implements the Bernoulli Naïve Bayes classifier to filter Reddit posts. Reddit is a social media form where users comment on content associated with different subreddits. For this project, each comment is associated with one of four subreddits: London, Montreal, Paris, and Toronto. The performance and accuracy were then compared to the Decision Tree and Support Vector Machines (SVM) methods using k-fold validation.

1.1 Classifiers

Decision Tree classification is a machine learning tool that uses values of each feature to recursively split the dataset to a point where all data points for the same class are grouped. SVM is another classifier that separates the training data by maximizing the margins between labels [1]. The python package *scikits.learn* was used to implement the Decision Tree and SVM classification. These two machine learning tools were selected since they are known to work well with text classification, which is compared to test the performance of the Bernoulli Naïve Bayes classifier. The hyper-parameters for the Decision Tree and SVM classification were selected for the highest k-fold accuracy but were not as good as the Bernoulli Naïve Bayes classifier.

There are no hyper-parameters associated with the Naïve Bayes classifier [2], however, there are preprocessing strategies applied to the features to ensure the best possible accuracy. Some of these strategies include implementing stop-words, lemmatization, n-grams, and language detection. Bernoulli Naïve Bayes only accepts binary data, which gave an accuracy of 77.5% even after preprocessing the data. Other Naïve Bayes approaches were implemented, including Laplace smoothing. By including these methods, acceptable accuracy was achieved.

1.2 Assumptions

Assumptions are required to properly build a model and train the data efficiently. The main assumption for this project was that the features, x_j are conditionally independent given the label y . This means that it is assumed that a feature presented in a class is unrelated to any other feature [3]. In mathematical terms,

$$P(x_j|y) = P(x_j|y, x_k). \quad (1)$$

2 Dataset Analysis

The datasets given for the project are user text from the Reddit website. The text is associated with one of four subreddits: London, Montreal, Paris, or Toronto. The features represent the words available in the text, while the labels are the associated subreddit. Two datasets are provided, one containing training data and one with test data. The training set has 718 samples that are relatively equally distributed amongst the four classes, and the test set has 279 samples.

2.1 Vectorization

CountVectorizer is a function within the sklearn Python library that tokenizes data. This package was used to split the words of the Reddit comments, however, some of the default settings were modified. The modifications were the removal of all accents, numbers, and punctuation. All words were also converted to lowercase, to ensure the same words are only accounted for once. The samples were then turned into binary features, where 1 represents the existence of a word and 0 represents it is non-existent in the sample. This method is effective for showing the presence of a word but lacks the ability to show its frequency in the sample.

2.2 Language Detection

Since the primary language in Toronto and London is English, while the primary language in Montreal and Paris is French, the language of the Reddit post was considered for the project. The langdetect package in Python is able to identify the language used in the post [4]. Two additional features were added to each post, one identifying the language as English, and the other as French.

2.3 Stop Words

Stop-words are common words that are likely to appear in multiple subreddits, despite their label. These words do not add any value to classifying the comments and therefore are removed from the feature list. In Python, there exist packages for common words in both English and French. Some examples of these words for both languages would include “and”, “the”, “he”, “et”, “le”, and “lui”. These words were removed from the text classification search, therefore putting more emphasis on more important words that will help properly identify the comment’s label.

2.4 Lemmatization and Stemmatization

Languages consist of several words that are derived from each other, for example, “change”, “changing”, and “changed”. Lemmatization and stemmatization are techniques that group together different forms of the same word. By combining the words together, the number of features is reduced, and a more efficient analysis is completed. Lemmatization keeps the root word of a similar word, which is more accurate for text classification. Stemmatization only takes the root of a word, without considering if the root is the readable word. For the previous example, lemmatization would group the words together as “change”, while stemmatization would group them together as “chang”.

2.5 N-Grams

N-grams is a method that uses a series of words instead of each individual word. The sequence of words will have a length n, which was limited in length to prevent a series of too many words. N-grams are important because some words could appear in the same order in multiple samples. For example, if you consider the trigram “The Three Musketeers”, these words would often appear in series which could be useful, whereas the term “three” alone may not be useful for classification.

2.6 Feature Selection

The number of features is reduced to the most frequent words observed throughout all posts. This is an important step because keeping too many features would result in over-fitting, but not enough features would result in insufficient classification. To reduce the number of features, the vector containing the binary values is reduced to the most frequent words observed throughout all posts. The number of words to keep was altered and validated using k-fold validation. This process is done using the ‘feature_selection’ class from *sklearn*. It selects the best n features by comparing their ANOVA F-value [5].

3 Proposed Approach

3.1 Considered Models

3.1.1 Bernoulli Naïve Bayes Classification

Naïve Bayes is a common text classification method because it is fast and easy to implement [3]. It is a probabilistic classifier, which means when given an input it will predict the probability of each class occurring. The Bernoulli Naïve Bayes classifier is a variant of the standard Naïve Bayes, where only discrete data is used and the features are provided in binary form. The assumptions for the classifier are specified in Section 1.2. The probability, $P(y = k|x)$, with k distinct output values is calculated for each class and the class with the highest probability is selected. The formula used to find the probability is

$$\begin{aligned} P(y = k|x) &\propto \delta_k(x) = \log P(y = k)P(x|y = k) \\ &= \log \theta_k + \sum_{j=1}^m (x_j * \log \theta_{j,k} + (1 - x_j) * \log(1 - \theta_{j,k})). \end{aligned} \quad (2)$$
$$(3)$$

From Laplace smoothing,

$$\theta_k = P(y = k) = \frac{\text{number of instances where}(y = k) + 1}{\text{total number of samples} + 2} \quad (4)$$

$$\theta_{j,k} = P(x_j = 1|y = k) = \frac{\text{number of instances where}(x_j = 1)\text{and}(y = k) + 1}{\text{total number of samples with}(y = k) + 2}. \quad (5)$$

The following decision rule assigns the class,

$$\text{Output} = \arg \max_x \delta_k(x). \quad (6)$$

3.1.2 Decision Tree

The Decision Tree classification was selected from the *scikit.learn* package to classify the data. Decision Trees classify the data by recursively splitting the data into different features. This process is repeated until all data points for the same class are grouped. The Decision Tree classifier was selected as an additional classifier because it requires little data preprocessing and is simple to implement. Three hyperparameters are considered:

- Tree depth: Maximum depth of the tree. Useful to reduce over-fitting. (Values considered: 50, 100, 500, 1000)
- Minimum sample split: Minimum number of samples in a node to split (Values considered: 2, 5, 10)
- Criterion: Which function to use to measure the quality of a feature. Two options are tested: *entropy* and *gini*.

3.1.3 SVM Classifier

A second classification method was selected to better compare the results. SVM classifier was selected as one of the additional *scikit.learn* packages. As discussed in Section 1, SVM classifier is a supervised learning method that classifies the data by maximizing the margins between labels. SVM was selected because it works well for high-dimensional spaces, which is relevant to text classification since each word represents a feature [6]. SVM classification is a common machine-learning method for text classification. The strength of regularization is the only hyperparameter considered. Values from 1 to $1e - 4$ will be tested. Lower values mean a stronger regularization and thus, less over-fitting.

3.2 Model Selection

The best model will be selected using their K-fold cross-validation accuracy and bias. Since a training and test dataset are provided, all the 718 training samples are used in training. K-fold cross-validation divides the dataset into k subsets, also known as folds. The data is trained and evaluated k times, using a different validation set each time. The average prediction accuracy over all k folds is computed. The model is then trained on the full dataset and its accuracy is measured. The difference between this accuracy and 100% is known as the bias.

To test all possible combinations of hyper-parameters with all the pre-processing options, an automated loop was implemented.

4 Results

The best models obtained for each classification method are presented in Table 1. To find the best combinations of hyper-parameters and pre-processing, an iterative process was used. The quality of the models was compared based on their 5-fold cross-validation accuracy and bias (1 - *training accuracy*).

First, the different pre-processing approaches were tested on the Bernoulli Naïve Bayes classifier. The best feature space for this model was found by following these steps:

1. Vary the number of features
2. Test different N-grams
3. Add language identification to features
4. Lemmatize the vocabulary

At each step, the best parameters are selected and kept constant for the next step. The best combinations of pre-processing were then used to train the other two models: SVM and decision trees. Note that when some combinations of pre-processing gave similar accuracies for Naïve Bayes, both were tested.

For SVM, a linear kernel with regularization was used. The impact of the regularization strength was analyzed. The effects of the criterion, tree depth, and minimum samples to split were tested for decision trees. For both models, the count of words was used instead of their binary value.

Also, since lemmatization did not improve the accuracy, and lemmatization of the dataset is computationally demanding, this option was not tested on the other model.

Table 1: Summarized results of the iterative process to find the best model

(a) Bernoulli Naïve Bayes			(b) Linear SVM			(c) Decision Tree		
Test	CV [%]	Bias [%]	Test	CV [%]	Bias [%]	Test	CV [%]	Bias [%]
Step 1 — N Features			Step 1 — Regularization			Step 1 — Criterion		
500	69.7	22.81	$C = 1$	71.3	0.0	entropy	55.6	0.0
1000	71.5	18.36	$C = 0.1$	74.0	0.0	gini	56.1	0.0
2000	74.8	11.54	$C = 0.05$	72.9	0.6	Step 2 — Max Depth		
3000	67.5	13.21	$C=0.01$	74.1	5.0	50	57.2	9.6
4000	62.9	13.21	$C = 0.005$	73.3	8.2	100	56.1	0
Step 2 — N-Grams			$C = 0.001$	68.4	17.5	500	55.6	0
2-grams	77.5	11.8	$C = 1e - 4$	65.7	25.0	1000	57.3	0
3-grams	74.3	12.8	Step 2 — N-Grams			Step 3 — Min Samples Split		
4-grams	73.2	15.3	2-grams	74.1	5.0	2	57.2	9.6
5-grams	70.8	18.2	3-grams	74.4	5.8	5	55.9	10.4
Step 3 — Language			Step 3 — Language			10	53.1	13.2
Without	77.5	11.8	Without	72.3	5.8			
With	76.1	13.0	With	72.2	8.1			
Step 4 — Lemmatization								
Without	77.5	11.8						
With	76.8	11.8						

From these results, it is now possible to select the model that will be used for the Kaggle competition. When selecting the model, it is important to consider the CV accuracy as well as the bias. A low bias i.e. a train accuracy of almost 100%, means that the model is overfitted. Based on this the following two models were chosen:

- Bernoulli Naïve Bayes, 2000 features, 2-grams and no language identification
- SVM, 2000 features, 3-grams with language identification

The test accuracies obtained are presented in Table 2.

Table 2: Best Models' Test Accuracies

Model	Test Accuracy [%]
Naïve Bayes	65.1
SVM	71.1

5 Discussion and Conclusion

Based on the results, SVM is the best model for the subreddit classification. The model's parameter and dataset pre-processing are reported below:

- N-grams: 2
- Feature selection: 2000 selected based on their ANOVA F-value
- Language identification: Yes
- Regularization: $C = 0.01$

This finding is surprising since, based on the cross-validation accuracy, it was expected that the Naïve Bayes classifier would perform better. However, its test accuracy is more than 10% lower while the one for SVM remained similar.

Processing the posts was one of the more delicate steps of the project. Because the posts were in two different languages, the processing of the text was complicated. To try to account for the different languages, an additional language feature was added. It was observed that the majority of the comments from the London or Toronto subreddits were English, and the majority of comments from the Montreal and Paris subreddits were French. An additional feature was added to represent the language as English or French. Although this additional feature had some advantages, there were issues with identifying posts from the minority language. For example, if there was an English post in the Montreal subreddit it would be classified as London or Toronto, despite some obvious indicators that it was from Montreal.

It is also interesting to analyze the feature probabilities of the Naïve Bayes classifier. As part of the training, the probability of occurrence of a certain feature for a given class is estimated by $P(x_j = 1|y = k) = \theta_{j,k}$. The words associated with the highest values of $\theta_{j,k}$ for the four classes are reported below.

Table 3: Most probable features for a given class

London	Montreal	Paris	Toronto
like	people	paris	people
london	montreal	plus	like
people	would	ca	one
get	get	si	toronto
also	like	tout	would

It is possible to note that some words, like “people“, are present in more than one class and thus, provide little information to the classification. A way to improve the accuracy of the Naïve Bayes classifier would be to weigh the words according to the information it provides about the class. This is the idea behind class-specific feature weighting. Class-specific weighting assigns a different weight to the feature depending on its relevance in the respective class, as opposed to the same weight across all classes [7]. This approach would benefit the classification of the project since there are certain words that have more relevance for certain classes. For instance, a word like “eiffel“ that is present only a few times (5) in the training dataset would be given a high weight since it was only observed when the class was “Paris“. By adding more weight to this word specifically for this class, better accuracy could be achieved. Ruan *et al.* [7] propose multiple feature weighting approaches that could be applied to text classification in this project.

Additionally, there are important characters such as \$, £, and € that are useful for classification. In addition to adding class-specific deep feature weighting to words, weight could be added to these features as well.

Another method for adding weight is term frequency x inverse document frequency (TF-IDF). TF-IDF adds weight to words that are more important in the text, meaning the word appears many times in the document and is a relatively rare word [8]. Since there are multiple classes in the sample, it would be interesting to add an additional component to the TF-IDF method to account for the frequency in a class compared to the overall dataset.

To conclude, two models with similar cross-validation performance were found: one using linear SVM and one using Bernoulli Naïve Bayes. However, finding the pre-processing that gave the best results was quite a challenging task that required a lot of intuition and time. For future work, it would be interesting to see if implementing class-specific feature weighting could reduce the dependence of the performances on this step.

6 Statement of Contribution

Charles created the Naïve Bayes algorithm and the k-fold validation code. He also implemented the automated process that compares the performances of the Naïve Bayes algorithm with the other two machine learning. Ashley created the code for preprocessing the features and wrote up the report.

References

- [1] A. Narges, “Ecse 551 - machine learning for engineers: Lecture 11 – regularization, decision trees,” 2023.
- [2] A. Narges, “Ecse 551 - machine learning for engineers: Lecture 8 – naive bayes,” 2023.
- [3] J. D. Rennie, L. Shih, J. Teevan, and D. R. Karger, “Tackling the poor assumptions of naive bayes text classifiers,” in *Proceedings of the 20th international conference on machine learning (ICML-03)*, 2003, pp. 616–623.
- [4] N. Shuyo, *Language detection library for java*, 2010. [Online]. Available: <http://code.google.com/p/language-detection/>.
- [5] [Online]. Available: https://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.f_classif.html#sklearn.feature_selection.f_classif.

- [6] S. Yue, P. Li, and P. Hao, “Svm classification: Its contents and challenges,” *Applied Mathematics-A Journal of Chinese Universities*, vol. 18, pp. 332–342, 2003.
- [7] S. Ruan, H. Li, C. Li, and K. Song, “Class-specific deep feature weighting for naïve bayes text classifiers,” *IEEE Access*, vol. 8, pp. 20 151–20 159, 2020.
- [8] A. Narges, “Ecse 551 - machine learning for engineers: Lecture 13 – decision trees (cont’d), feature construction,” 2023.

7 Appendix

7.1 MP2.ipynb

MP2

November 19, 2023

1 ECSE-551 Mini Project 2

Authors: * Ashley Meagher (260822930) * Charles Sirois (261158513)

```
[250]: # To specify where to load the data
in_colab = True
folder_path = 'drive/MyDrive/Colab Notebooks/ECSE 551_MP2'

%load_ext autoreload
%autoreload 2

# Our functions and classes
if in_colab:
    from google.colab import drive
    from google.colab import data_table
    drive.mount('/content/drive')

    data_table.enable_dataframe_formatter() # For interactive df viz

    import sys
    sys.path.insert(0, folder_path)

# SK Learn models
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import LinearSVC

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time
import itertools
import datetime

import nltk
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('averaged_perceptron_tagger')
```

```

# Install required packages
!pip install unicode # To remove accents
!pip install langid # To identify text's language

# Import our classes and functions from the other files
from NaiveBayes import NaiveBayes
from cross_val_score import cross_val_score
from data_processing import Data, Format_data

```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Requirement already satisfied: unicode in /usr/local/lib/python3.10/dist-packages (1.3.7)

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /root/nltk_data...
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!

```

Requirement already satisfied: langid in /usr/local/lib/python3.10/dist-packages (1.1.6)

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from langid) (1.23.5)

1.1 Data Analysis

1.1.1 Load the data

```

[251]: print(f"Loading data files... ", end='')
filenames = [folder_path + "/data/train_utf8.csv", folder_path + "/data/
↳ test_utf8.csv"]
words_dataset = Data(train_file=filenames[0], test_file=filenames[1])
print(f'Done')

```

Loading data files... Done

1.1.2 Data properties

```
[252]: print(f'Training dataset size: {words_dataset.train_data.shape[0]}')
print(f'Test dataset size: {words_dataset.test_data.shape[0]}')

classes, classes_count = np.unique(words_dataset.train_data['label'],
    ↪return_counts=True)
print(f'Classes: ')
for cls, cls_count in zip(classes, classes_count):
    print(f'\t-{cls}: {cls_count}')
```

Training dataset size: 719

Test dataset size: 279

Classes:

- London: 180
- Montreal: 180
- Paris: 180
- Toronto: 179

1.2 Models Performances

1.2.1 Functions

Functions to compute the cross-validation score of the different combinations of model hyperparameters and datasets

```
[253]: def create_datasets(ds_options_dict):
    """
    To create a list with all the combinations of options in the dict
    """
    print(f"Processing input data...")
    keys, values = zip(*ds_options_dict.items())
    ds_options_list = [dict(zip(keys, v)) for v in itertools.product(*values)]

    ds_list = []
    for idx, each_ds in enumerate(ds_options_list):
        each_ds['dataset_name'] = f'DS {idx}'
        ds_list.append(Format_data(words_dataset, **each_ds))

    print(f'\nDone')
    return ds_list

def find_ds_from_name(ds_name, ds_list) -> Format_data:
    """
    To return the dataset with the corresponding name in the `ds_list`
    """
    ds = next((ds for ds in ds_list if ds.name == ds_name), None)

    if ds is None:
```

```

        raise ValueError(f"Dataset {ds_name} not found in `ds_list`")

    return ds

def compute_models_cv_acc(model_dict, ds_list):
    """
    To compute the cv score for all the combinations of model_dict and ds_list
    """
    results_df = pd.DataFrame()

    # Cross-Validation
    n_fold = 5

    start_time = time.time()
    print(f"----- Training all models -----")
    for model_name, model_info in model_dict.items():
        model = model_info["model"]
        base_params = model_info["base_params"]
        cv_params = model_info["cv_params"]

        print(f"\nModel : {model_name}")
        model_start = time.time()
        for ds_idx, each_dataset in enumerate(ds_list):
            # Check if it already has been ran
            ds_start = time.time()
            dataset_name = each_dataset.name
            print(f"\tDataset [{ds_idx+1}/{len(ds_list)}]: {dataset_name}")

            X_train = each_dataset.X
            y_train = each_dataset.Y

            # Cross_validation
            cv_results = cross_val_score(
                model,
                X_train,
                y_train,
                cv=n_fold,
                base_params=base_params,
                cv_params=cv_params,
                results_df=results_df,
                ds_name=dataset_name,
            )

            if cv_results.empty:
                print(f'... Model already trained')
                continue

```

```

# Print best combination
best_row = cv_results.iloc[cv_results['Score'].idxmax()]
compute_time = time.time() - ds_start
print(
    f"\tBest CV Score : {np.round(best_row['Score']*100)}% (Acc: {np.
↪round(best_row['Acc']*100)}) "
    f"[{compute_time} sec]\n"
)

# Add information to series
ds_params = each_dataset.get_params()

for key, value in ds_params.items():
    if isinstance(value, tuple):
        value = str(value)
    cv_results[key] = value

# cv_results = pd.concat([cv_results, pd.(ds_params).T],
↪ignore_index=True)
cv_results['Model name'] = model_name
cv_results['Dataset'] = dataset_name
cv_results['Compute time'] = compute_time

results_df = pd.concat([results_df, cv_results], ignore_index=True,
↪axis=0)

print(f'Model trained in {time.time() - model_start} sec')

print(f"\nTraining completed ({time.time() - start_time} sec)\n")

results_df['Score'] = (results_df['Score']*100).apply(np.round, decimals=2)
results_df['Bias'] = ((1 - results_df['Acc'])*100).apply(np.round, decimals=2)

results_df = results_df[
    [
        'Model name',
        'Score',
        'Bias',
        'Acc',
        'Dataset',
        'Params',
        'Compute time',
        'Model',
        'n_gram',
        'feat_type',
        'lemmatized',
        'lang',
    ]

```

```

        'standardized',
        'rm_accents',
        'feat_select',
        'n_feat',
    ]
]

results_df = results_df.sort_values(by=['Score'], ascending=False)

return results_df

def create_pred_ds(model_idx, results_df, ds_list, save_path):
    """ To create the prediction csv file for the model corresponding to
    ↪ model_idx
    The file is saved under save_path
    """
    my_model_info = results_df.loc[model_idx]
    print(f'Model chosen: ')
    print(my_model_info)

    print(f"Predicting test data using this model...")
    my_model = my_model_info['Model']
    ds = find_ds_from_name(my_model_info['Dataset'], ds_list)

    y_test = my_model.predict(ds.X_test)
    pred_df = pd.DataFrame(y_test, columns=['subreddit'])
    pred_df.index.name = 'id'

    pred_df.to_csv(save_path)
    print(f'Predictions saved to {save_path}')

```

1.2.2 Bernoulli Naive Bayes

Parameters evaluation

```

[254]: nb_ds_options = {
        'max_feat': [None],
        'lang_id': [False, True], # [False, True],
        'feature_type': ['Bin'],
        'n_gram': [(1, 1), (1, 2), (1, 3), (1, 4), (1, 5)],
        'lemmatize': [False],
        'feat_select': ['F_CL'],
        'n_feat_select': [500, 1000, 2000, 3000, 4000],
    }
nb_ds_list = create_datasets(nb_ds_options)

```

```
Processing input data...
    Processing of: DS 49...
Done
```

```
[255]: # Separate datasets to test lemmatization
nb2_ds_options = {
    'max_feat': [None],
    'lang_id': [False], # [False, True],
    'feature_type': ['Bin'],
    'n_gram': [(1, 2)],
    'lemmatize': [True, False],
    'feat_select': ['F_CL'],
    'n_feat_select': [2000],
}
nb2_ds_list = create_datasets(nb2_ds_options)
```

```
Processing input data...
    Processing of: DS 1...
Done
```

```
[256]: nb_model_dict = {}
nb_model_dict["My Bernouilli NB"] = {
    "model": NaiveBayes,
    'base_params': {'laplace_smoothing': True, 'verbose': False},
    'cv_params': None,
}
```

```
[257]: nb_df = compute_models_cv_acc(nb_model_dict, nb_ds_list)
nb_results = nb_df[['Score', 'Bias', 'Params', 'n_gram', 'lang', 'n_feat']]
```

```
----- Training all models -----
```

```
Model : My Bernouilli NB
Dataset [1/50]: DS 0
Combination 1/1 Best CV Score : 71.0% (Acc: 77.0) [4.215232849121094
sec]

Dataset [2/50]: DS 1
Combination 1/1 Best CV Score : 72.0% (Acc: 82.0) [8.167191982269287
sec]

Dataset [3/50]: DS 2
Combination 1/1 Best CV Score : 75.0% (Acc: 88.0) [20.4975848197937 sec]

Dataset [4/50]: DS 3
Combination 1/1 Best CV Score : 68.0% (Acc: 87.0) [26.63149333000183
sec]

Dataset [5/50]: DS 4
```

Combination 1/1 Best CV Score : 62.0% (Acc: 83.0) [37.83934044837952
sec]

Dataset [6/50]: DS 5
Combination 1/1 Best CV Score : 70.0% (Acc: 77.0) [5.771634101867676
sec]

Dataset [7/50]: DS 6
Combination 1/1 Best CV Score : 71.0% (Acc: 83.0) [7.290061712265015
sec]

Dataset [8/50]: DS 7
Combination 1/1 Best CV Score : 78.0% (Acc: 88.0) [18.11064648628235
sec]

Dataset [9/50]: DS 8
Combination 1/1 Best CV Score : 74.0% (Acc: 91.0) [29.609700202941895
sec]

Dataset [10/50]: DS 9
Combination 1/1 Best CV Score : 68.0% (Acc: 88.0) [36.51623868942261
sec]

Dataset [11/50]: DS 10
Combination 1/1 Best CV Score : 65.0% (Acc: 71.0) [4.02064323425293 sec]

Dataset [12/50]: DS 11
Combination 1/1 Best CV Score : 70.0% (Acc: 81.0) [10.331190824508667
sec]

Dataset [13/50]: DS 12
Combination 1/1 Best CV Score : 75.0% (Acc: 87.0) [18.814451694488525
sec]

Dataset [14/50]: DS 13
Combination 1/1 Best CV Score : 73.0% (Acc: 90.0) [29.117358684539795
sec]

Dataset [15/50]: DS 14
Combination 1/1 Best CV Score : 69.0% (Acc: 89.0) [37.6209557056427 sec]

Dataset [16/50]: DS 15
Combination 1/1 Best CV Score : 60.0% (Acc: 66.0) [3.796255588531494
sec]

Dataset [17/50]: DS 16
Combination 1/1 Best CV Score : 68.0% (Acc: 76.0) [11.544331073760986
sec]

Dataset [18/50]: DS 17
 Combination 1/1 Best CV Score : 72.0% (Acc: 85.0) [17.99778699874878
 sec]

Dataset [19/50]: DS 18
 Combination 1/1 Best CV Score : 70.0% (Acc: 89.0) [28.106534481048584
 sec]

Dataset [20/50]: DS 19
 Combination 1/1 Best CV Score : 69.0% (Acc: 89.0) [36.949981927871704
 sec]

Dataset [21/50]: DS 20
 Combination 1/1 Best CV Score : 57.0% (Acc: 62.0) [3.6437978744506836
 sec]

Dataset [22/50]: DS 21
 Combination 1/1 Best CV Score : 64.0% (Acc: 74.0) [9.618980169296265
 sec]

Dataset [23/50]: DS 22
 Combination 1/1 Best CV Score : 69.0% (Acc: 82.0) [18.49955105781555
 sec]

Dataset [24/50]: DS 23
 Combination 1/1 Best CV Score : 69.0% (Acc: 86.0) [26.71236515045166
 sec]

Dataset [25/50]: DS 24
 Combination 1/1 Best CV Score : 68.0% (Acc: 88.0) [36.23162865638733
 sec]

Dataset [26/50]: DS 25
 Combination 1/1 Best CV Score : 69.0% (Acc: 75.0) [6.965734243392944
 sec]

Dataset [27/50]: DS 26
 Combination 1/1 Best CV Score : 70.0% (Acc: 81.0) [7.405640363693237
 sec]

Dataset [28/50]: DS 27
 Combination 1/1 Best CV Score : 74.0% (Acc: 87.0) [17.702255725860596
 sec]

Dataset [29/50]: DS 28
 Combination 1/1 Best CV Score : 67.0% (Acc: 85.0) [28.42513680458069
 sec]

Dataset [30/50]: DS 29
 Combination 1/1 Best CV Score : 63.0% (Acc: 82.0) [36.266666412353516
 sec]

Dataset [31/50]: DS 30
 Combination 1/1 Best CV Score : 69.0% (Acc: 75.0) [3.7581331729888916
 sec]

Dataset [32/50]: DS 31
 Combination 1/1 Best CV Score : 73.0% (Acc: 83.0) [10.553314924240112
 sec]

Dataset [33/50]: DS 32
 Combination 1/1 Best CV Score : 75.0% (Acc: 87.0) [17.867765426635742
 sec]

Dataset [34/50]: DS 33
 Combination 1/1 Best CV Score : 74.0% (Acc: 88.0) [26.27276086807251
 sec]

Dataset [35/50]: DS 34
 Combination 1/1 Best CV Score : 67.0% (Acc: 87.0) [37.33134150505066
 sec]

Dataset [36/50]: DS 35
 Combination 1/1 Best CV Score : 62.0% (Acc: 69.0) [4.442660331726074
 sec]

Dataset [37/50]: DS 36
 Combination 1/1 Best CV Score : 70.0% (Acc: 79.0) [7.428954124450684
 sec]

Dataset [38/50]: DS 37
 Combination 1/1 Best CV Score : 75.0% (Acc: 86.0) [18.159637451171875
 sec]

Dataset [39/50]: DS 38
 Combination 1/1 Best CV Score : 72.0% (Acc: 88.0) [28.107608318328857
 sec]

Dataset [40/50]: DS 39
 Combination 1/1 Best CV Score : 69.0% (Acc: 87.0) [35.64508605003357
 sec]

Dataset [41/50]: DS 40
 Combination 1/1 Best CV Score : 59.0% (Acc: 64.0) [4.095541715621948
 sec]


```
Dataset [42/50]: DS 41
Combination 1/1 Best CV Score : 66.0% (Acc: 75.0) [10.126330375671387
sec]
```

```
Dataset [43/50]: DS 42
Combination 1/1 Best CV Score : 72.0% (Acc: 84.0) [17.864989519119263
sec]
```

```
Dataset [44/50]: DS 43
Combination 1/1 Best CV Score : 70.0% (Acc: 87.0) [27.56903648376465
sec]
```

```
Dataset [45/50]: DS 44
Combination 1/1 Best CV Score : 69.0% (Acc: 88.0) [37.22462606430054
sec]
```

```
Dataset [46/50]: DS 45
Combination 1/1 Best CV Score : 56.0% (Acc: 60.0) [3.685528516769409
sec]
```

```
Dataset [47/50]: DS 46
Combination 1/1 Best CV Score : 64.0% (Acc: 72.0) [9.841209650039673
sec]
```

```
Dataset [48/50]: DS 47
Combination 1/1 Best CV Score : 69.0% (Acc: 81.0) [18.502774715423584
sec]
```

```
Dataset [49/50]: DS 48
Combination 1/1 Best CV Score : 67.0% (Acc: 84.0) [25.65246319770813
sec]
```

```
Dataset [50/50]: DS 49
Combination 1/1 Best CV Score : 69.0% (Acc: 87.0) [35.62777662277222
sec]
```

Model trained in 964.5682625770569 sec

Training completed (964.5693309307098 sec)

```
[258]: nb2_df = compute_models_cv_acc(nb_model_dict, nb2_ds_list)
nb2_results = nb2_df[['Score', 'Bias', 'Params', 'n_gram', 'lang', 'n_feat',
↳ 'lemmatized']]
```

----- Training all models -----

```

Model : My Bernouilli NB
      Dataset [1/2]: DS 0
      Combination 1/1 Best CV Score : 75.0% (Acc: 88.0) [18.69008207321167
sec]

      Dataset [2/2]: DS 1
      Combination 1/1 Best CV Score : 79.0% (Acc: 88.0) [19.127805709838867
sec]

Model trained in 37.849053144454956 sec

Training completed (37.85080671310425 sec)

```

Step 1 - Effect of feature selection

```

[259]: nb_results_step1 = nb_results[(nb_results['n_gram'] == '(1, 1)') &
      ↪(nb_results['lang'] == False)]
      nb_results_step1.sort_values(by=['n_feat'], ascending=True)

```

```

[259]:
   Score  Bias Params  n_gram  lang  n_feat
0  71.36  22.81      {}  (1, 1)  False    500
1  71.63  18.36      {}  (1, 1)  False   1000
2  75.11  11.54      {}  (1, 1)  False   2000
3  67.59  13.21      {}  (1, 1)  False   3000
4  61.89  16.97      {}  (1, 1)  False   4000

```

Step 2 - N-grams

```

[260]: nb_results_step2 = nb_results[(nb_results['n_feat'] == 2000) &
      ↪(nb_results['lang'] == False)]
      nb_results_step2.sort_values(by=['n_gram'], ascending=True)

```

```

[260]:
   Score  Bias Params  n_gram  lang  n_feat
2  75.11  11.54      {}  (1, 1)  False   2000
7  78.17  11.82      {}  (1, 2)  False   2000
12 74.83  12.80      {}  (1, 3)  False   2000
17 72.05  15.30      {}  (1, 4)  False   2000
22 69.12  18.22      {}  (1, 5)  False   2000

```

Step 3 - Language Identification

```

[261]: nb_results_step3 = nb_results[(nb_results['n_feat'] == 2000) &
      ↪(nb_results['n_gram'] == '(1, 2)')]
      nb_results_step3

```

```

[261]:
   Score  Bias Params  n_gram  lang  n_feat
7  78.17  11.82      {}  (1, 2)  False   2000
32 75.25  12.93      {}  (1, 2)  True    2000

```

Step 4 - Effect of lemmatization

```
[262]: nb_results_step4 = nb2_results
       nb_results_step4
```

```
[262]:
```

	Score	Bias	Params	n_gram	lang	n_feat	lemmatized
1	78.73	11.82	{}	(1, 2)	False	2000	False
0	75.11	11.82	{}	(1, 2)	False	2000	True

1.2.3 SVC

```
[263]: svc_ds_options = {
       'max_feat': [None],
       'lang_id': [False, True], # [False, True],
       'feature_type': ['Count'],
       'n_gram': [(1, 2), (1, 3)],
       'lemmatize': [False],
       'feat_select': ['F_CL'],
       'n_feat_select': [2000],
       }
       svc_ds_list = create_datasets(svc_ds_options)

       svc_model_dict = {}
       svc_model_dict["SVC"] = {
           "model": LinearSVC,
           "base_params": {"random_state": 0},
           "cv_params": {"C": [0.0001, 0.001, 0.005, 0.01, 0.05, 0.1, 1]},
       }
```

Processing input data...

Processing of: DS 3...

Done

```
[264]: svc_df = compute_models_cv_acc(svc_model_dict, svc_ds_list)
       svc_results = svc_df[['Score', 'Bias', 'Params', 'n_gram', 'lang', 'n_feat']]
```

----- Training all models -----

Model : SVC

Dataset [1/4]: DS 0

Combination 7/7 Best CV Score : 74.0% (Acc: 95.0) [0.5722982883453369

sec]

Dataset [2/4]: DS 1

Combination 7/7 Best CV Score : 73.0% (Acc: 99.0) [0.7092363834381104

sec]

Dataset [3/4]: DS 2

Combination 7/7 Best CV Score : 74.0% (Acc: 99.0) [0.5993454456329346

sec]

Dataset [4/4]: DS 3

Combination 7/7 Best CV Score : 74.0% (Acc: 99.0) [1.3745067119598389

sec]

Model trained in 3.2752864360809326 sec

Training completed (3.2758843898773193 sec)

[265]: svc_results

[265]:	Score	Bias	Params	n_gram	lang	n_feat
3	74.13	5.01	{'C': 0.01}	(1, 2)	False	2000
4	73.99	0.56	{'C': 0.05}	(1, 2)	False	2000
18	73.85	0.70	{'C': 0.05}	(1, 2)	True	2000
25	73.85	0.83	{'C': 0.05}	(1, 3)	True	2000
2	73.57	8.21	{'C': 0.005}	(1, 2)	False	2000
5	73.44	0.00	{'C': 0.1}	(1, 2)	False	2000
11	73.16	0.70	{'C': 0.05}	(1, 3)	False	2000
19	72.60	0.00	{'C': 0.1}	(1, 2)	True	2000
12	72.47	0.14	{'C': 0.1}	(1, 3)	False	2000
17	72.47	7.93	{'C': 0.01}	(1, 2)	True	2000
9	72.47	8.21	{'C': 0.005}	(1, 3)	False	2000
10	72.32	5.84	{'C': 0.01}	(1, 3)	False	2000
26	72.18	0.14	{'C': 0.1}	(1, 3)	True	2000
24	71.90	8.07	{'C': 0.01}	(1, 3)	True	2000
16	70.93	12.52	{'C': 0.005}	(1, 2)	True	2000
27	70.52	0.00	{'C': 1}	(1, 3)	True	2000
20	70.51	0.00	{'C': 1}	(1, 2)	True	2000
13	69.54	0.00	{'C': 1}	(1, 3)	False	2000
23	69.12	12.38	{'C': 0.005}	(1, 3)	True	2000
6	68.99	0.00	{'C': 1}	(1, 2)	False	2000
8	68.01	17.39	{'C': 0.001}	(1, 3)	False	2000
15	67.31	22.67	{'C': 0.001}	(1, 2)	True	2000
1	66.20	17.52	{'C': 0.001}	(1, 2)	False	2000
7	65.49	25.31	{'C': 0.0001}	(1, 3)	False	2000
22	65.09	22.67	{'C': 0.001}	(1, 3)	True	2000
0	64.25	25.03	{'C': 0.0001}	(1, 2)	False	2000
14	59.66	31.29	{'C': 0.0001}	(1, 2)	True	2000
21	59.39	31.57	{'C': 0.0001}	(1, 3)	True	2000

Step 1 - Regularization

```
[266]: svc_results_step1 = svc_results[(svc_results['n_feat'] == 2000) &
    ↪(svc_results['n_gram'] == '(1, 2)') & (svc_results['lang'] == False)]
svc_results_step1
```

```
[266]:
```

	Score	Bias	Params	n_gram	lang	n_feat
3	74.13	5.01	{'C': 0.01}	(1, 2)	False	2000
4	73.99	0.56	{'C': 0.05}	(1, 2)	False	2000
2	73.57	8.21	{'C': 0.005}	(1, 2)	False	2000
5	73.44	0.00	{'C': 0.1}	(1, 2)	False	2000
6	68.99	0.00	{'C': 1}	(1, 2)	False	2000
1	66.20	17.52	{'C': 0.001}	(1, 2)	False	2000
0	64.25	25.03	{'C': 0.0001}	(1, 2)	False	2000

Step 2 - N-Grams

```
[267]: svc_results_step2 = svc_results[(svc_results['n_feat'] == 2000) &
    ↳(svc_results['Params'] == {'C': 0.01}) & (svc_results['lang'] == False)]
svc_results_step2
```

```
[267]:
```

	Score	Bias	Params	n_gram	lang	n_feat
3	74.13	5.01	{'C': 0.01}	(1, 2)	False	2000
10	72.32	5.84	{'C': 0.01}	(1, 3)	False	2000

Step 3 - Language

```
[268]: svc_results_step3 = svc_results[(svc_results['n_feat'] == 2000) &
    ↳(svc_results['Params'] == {'C': 0.01}) & (svc_results['n_gram'] == '(1, 3)')]
svc_results_step3
```

```
[268]:
```

	Score	Bias	Params	n_gram	lang	n_feat
10	72.32	5.84	{'C': 0.01}	(1, 3)	False	2000
24	71.90	8.07	{'C': 0.01}	(1, 3)	True	2000

1.2.4 Decision Tree

```
[269]: dt_ds_options = {
    'max_feat': [None],
    'lang_id': [False], # [False, True],
    'feature_type': ['Count'],
    'n_gram': [(1, 2)],
    'lemmatize': [False],
    'feat_select': ['F_CL'],
    'n_feat_select': [2000],
}
dt_ds_list = create_datasets(dt_ds_options)

dt_model_dict = {}
dt_model_dict["DT"] = {
    "model": DecisionTreeClassifier,
    "base_params": {"random_state": 0},
    "cv_params": {"criterion": ['gini', 'entropy'],
                  "max_depth": [50, 100, 500, 1000],
```

```

        "min_samples_split": [2, 5, 10]},
    }

```

Processing input data...

Processing of: DS 0...

Done

```

[270]: dt_df = compute_models_cv_acc(dt_model_dict, dt_ds_list)

dt_df['criterion'] = None
dt_df['max_depth'] = None
dt_df['min_samples_split'] = None

for idx, each_row in dt_df.iterrows():
    for key, val in each_row['Params'].items():
        dt_df.at[idx, key] = val

dt_results = dt_df[['Score', 'Bias', 'criterion', 'max_depth', '
↳ 'min_samples_split', 'n_gram', 'lang', 'n_feat']]

```

----- Training all models -----

Model : DT

Dataset [1/1]: DS 0

Combination 24/24

Best CV Score : 59.0% (Acc: 100.0)

[8.762712717056274 sec]

Model trained in 8.768966913223267 sec

Training completed (8.770277261734009 sec)

Step 1 - Criterion

```

[271]: dt_results_step1 = dt_results[(dt_results['n_feat'] == 2000) &
↳ (dt_results['max_depth'] == 100) & (dt_results['min_samples_split'] == 2)]
dt_results_step1

```

```

[271]:      Score  Bias criterion max_depth min_samples_split  n_gram  lang  n_feat
3    58.28   0.0      gini      100                2  (1, 2)  False   2000
15   54.39   0.0    entropy      100                2  (1, 2)  False   2000

```

Step 2 - Max Depth

```

[272]: dt_results_step2 = dt_results[(dt_results['n_feat'] == 2000) &
↳ (dt_results['criterion'] == 'gini') & (dt_results['min_samples_split'] == 2)]
dt_results_step2

```

```
[272]:
```

	Score	Bias	criterion	max_depth	min_samples_split	n_gram	lang	n_feat
9	58.55	0.0	gini	1000		2 (1, 2)	False	2000
3	58.28	0.0	gini	100		2 (1, 2)	False	2000
6	57.17	0.0	gini	500		2 (1, 2)	False	2000
0	55.91	9.6	gini	50		2 (1, 2)	False	2000

Step 3 - Min Samples Split

```
[273]: dt_results_step3 = dt_results[(dt_results['n_feat'] == 2000) &
↳(dt_results['criterion'] == 'gini') & (dt_results['max_depth'] == 50)]
dt_results_step3
```

```
[273]:
```

	Score	Bias	criterion	max_depth	min_samples_split	n_gram	lang	n_feat
1	56.47	10.43	gini	50		5 (1, 2)	False	2000
0	55.91	9.60	gini	50		2 (1, 2)	False	2000
2	51.87	13.21	gini	50		10 (1, 2)	False	2000

1.2.5 Final Model

```
[274]: best_nb_idx = 7
best_svc_idx = 24
```

```
[275]: create_pred_ds(best_nb_idx, nb_df, nb_ds_list, folder_path + '/'
↳NB_Final_prediction.csv')
```

Model chosen:

Model name	My Bernouilli NB
Score	78.17
Bias	11.82
Acc	0.88178
Dataset	DS 7
Params	{}
Compute time	18.110646
Model	<NaiveBayes.NaiveBayes object at 0x78dce2ed2d40>
n_gram	(1, 2)
feat_type	Bin
lemmatized	False
lang	False
standardized	False
rm_accents	True
feat_select	F_CL
n_feat	2000

Name: 7, dtype: object

Predicting test data using this model...

Predictions saved to drive/MyDrive/Colab Notebooks/ECSE

551_MP2/NB_Final_prediction.csv

```
[276]: create_pred_ds(best_svc_idx, svc_df, svc_ds_list, folder_path + '/'
↳SVC_Final_prediction.csv')
```

```
Model chosen:
Model name          SVC
Score               71.9
Bias                8.07
Acc                 0.919332
Dataset             DS 3
Params              {'C': 0.01}
Compute time        1.374507
Model               LinearSVC(C=0.01, random_state=0)
n_gram              (1, 3)
feat_type           Count
lemmatized          False
lang                True
standardized        False
rm_accents          True
feat_select         F_CL
n_feat              2000
Name: 24, dtype: object
Predicting test data using this model...
Predictions saved to drive/MyDrive/Colab Notebooks/ECSE
551_MP2/SVC_Final_prediction.csv
```

1.3 Features Analysis

Check the words that are most probably observed for a given class

```
[277]: nb_info = nb_df.loc[best_nb_idx]
nb_model = nb_info['Model']
nb_model_ds = find_ds_from_name(nb_info['Dataset'], nb_ds_list)
features_name = nb_model_ds.features_name

n_best_features = 10

df_dict = {}
for k, class_label in enumerate(nb_model._classes):
    feats_score = nb_model._thetas[k, 1:]
    names_scores = list(zip(features_name, feats_score))
    feat_scores_df = pd.DataFrame(data=names_scores, columns=['Feat_names',
↳'Score'])
    feat_scores_df = feat_scores_df.sort_values(by=['Score'], ascending=False).
↳reset_index(drop=True)
    df_dict[class_label] = feat_scores_df

combined_df = pd.concat(df_dict, axis=1)
# print(combined_df.head(n_feats).to_string())
```



```
combined_df.head(n_best_features)
```

```
[277]:
```

	London		Montreal		Paris		Toronto \	
	Feat_names	Score	Feat_names	Score	Feat_names	Score	Feat_names	
0	like	0.296703	people	0.225275	paris	0.318681	people	
1	london	0.280220	montreal	0.186813	plus	0.269231	like	
2	people	0.263736	would	0.175824	ca	0.263736	one	
3	one	0.252747	get	0.175824	si	0.203297	toronto	
4	get	0.225275	like	0.159341	tout	0.181319	would	
5	also	0.192308	one	0.153846	etre	0.164835	city	
6	really	0.159341	ca	0.142857	faire	0.142857	also	
7	would	0.153846	good	0.126374	quand	0.142857	get	
8	know	0.153846	go	0.120879	comme	0.131868	time	
9	see	0.142857	time	0.115385	fait	0.126374	new	

	Score
0	0.248619
1	0.243094
2	0.204420
3	0.198895
4	0.187845
5	0.171271
6	0.165746
7	0.160221
8	0.160221
9	0.149171

7.2 NaiveBayes.py

```
1 from typing import Literal
2 import numpy as np
3 from mpl_toolkits.mplot3d import Axes3D
4 import matplotlib.pyplot as plt
5 import pandas as pd
6 import time
7
8
9 def sigmoid(x):
10     return 1 / (1 + np.exp(-x))
11
12
13 def cost_function(X, y, w):
14     n = len(y)
15     h = sigmoid(np.dot(X, w))
16     J = (-1 / n) * (np.dot(y.T, np.log(h)) + np.dot((1 - y).T, np.log(1 - h)))
17     return J
18
19
20 class NaiveBayes:
21     """Naive Bayes class.
22
23     If k is different than 0, will find the best alpha, tol and reg_cst
24     using k-fold cross-validation.
25
26     Attributes:
27         thetas (ndarray): ndarray of shape (k, (1 + m))
28             Theta values for classification
29         class: Output class
30         n_features (int): Dimension of features (m). 0 if model not fitted.
31         n_iter (int): Number of iterations needed for fitting. 0 if not fitted.
32
33         X (ndarray): Training inputs (nxm)
34         y (ndarray): Training output (nx1)
35     """
36
37     def __init__(
38         self,
39         laplace_smoothing: bool = True,
40         verbose=False, # To print execution info
41     ) -> None:
42         self._n_features = 0
43
44         self._classes = None
45         self._n_class = 0
46         self._class_count = None
47         self._n_samples = 0
48
49         self.X = None # X dataset for training
50         self.y = None # Y dataset for training
51
52         self.laplace_smoothing = laplace_smoothing
53
54         self._log_class_prior = None
55         self._feat_log_proba = None
56         self._feat_log_proba = None
57
58         self.results = None # Results dataframe
59         self._comp_time = None
60
61         self._verbose = verbose
62
63     def fit(self, X, y):
```

```

64 """
65 Fit the model according to the given training data.
66
67 Parameters:
68     X (ndarray) : shape (n_samples, n_features)
69                 Training vector.
70
71     y (ndarray) : shape (n_samples,)
72                 Expected output vector
73
74     w (ndarray, optional): shape (n_features, )
75     T o give an initial guess
76
77 Returns:
78     self
79     Model with weights fitted to training dataset
80 """
81 self.X = X
82 self.y = y
83
84 # Check dataset sizes
85 self._n_samples, self._n_features = X.shape
86
87 if X.shape[0] != y.shape[0]:
88     raise ValueError(f"Mismatch between the size of the input ({X.shape[0]})
and outputs ({y.shape[0]})")
89
90 # Number of class
91 self._classes, self._class_count = np.unique(self.y, return_counts=True)
92 self._n_class = len(self._classes)
93
94 if self._verbose:
95     print(f"Fitting Naive Bayes model for the dataset")
96     print(f"\t# Features: {self._n_features}, classes: {self._classes}, #
Samples: {self._n_samples}")
97
98 self._train_model()
99
100 return self
101
102 def predict(self, X):
103     """To predict the output of samples.
104
105     For each class, computes:
106     $$ \delta_k = \log \{ [\theta_k \sum_{j=1}^m \theta_{j,k}^{x_j} (1 - \theta_{j,k})^{1-x_j}] \} $$
107
108     Classify the output as:
109     $$ \text{Output} = \arg\max_k \delta_k(x) $$
110
111     Args:
112     X (ndarray): Inputs sample to be predicted. Size (n x m)
113
114     Raises:
115     ValueError: If model is not trained
116
117     Returns:
118     ndarray: Predicted output of each data point (n x 1)
119     """
120     if self._thetas is None:
121         raise ValueError(f"Model is not trained.")
122
123     log_class_prior = np.log(self._thetas[:, 0]) # log( P(Y=k) ) shape: 1 x k
124     feat_log_proba = np.log(self._thetas[:, 1:]) # log( P(x_j=1 | Y=k) ) shape:
k x m

```

```

125     feat_log_neg_proba = np.log(1 - self._thetas[:, 1::]) # log( 1 - P(x_j=1 | Y=
k) )   shape: kxm
126
127     n_samples = X.shape[0]
128
129     self._joint_log_likelihood = np.zeros((n_samples, self._n_class))
130
131     self._joint_log_likelihood = X @ (feat_log_proba - feat_log_neg_proba).T +
feat_log_neg_proba.sum(axis=1)
132
133     # Add class log priors
134     self._joint_log_likelihood += log_class_prior
135
136     # Predictions
137     predictions_idx = np.argmax(self._joint_log_likelihood, axis=1)
138     predictions = self._classes[predictions_idx]
139
140     return predictions
141
142 def score(self, X, y):
143     """To compute the accuracy of the model
144
145     Args:
146         X (ndarray): Test samples
147         y (ndarray): True class of X
148
149     Returns:
150         float: Accuracy of the model over test samples
151     """
152     y_pred = self.predict(X)
153
154     accuracy = (y == y_pred).mean()
155
156     return accuracy
157
158 def _train_model(self):
159     """Compute the theta needed to estimate the probabilities
160
161     For each class:
162         $$ \theta_k = P(Y=k) = (\text{\# samples where } Y=k) / (\text{\# samples}) $$
163         $$ \theta_{\{j, k\}} = P(x_j=1 | Y=k) = (\text{\# samples where } x_j=1 \text{ and } Y=k) / (\text{\#
samples where } Y=k) $$
164
165     Stores the values in a np.array of shape k x (1 + m)
166         -> \theta_{k} = thetas[k, 0], k=0, ... n_class (prior of class k)
167         -> \theta_{\{j, k\}} = thetas[k, j], j=1, ..., m
168     """
169     self._thetas = np.zeros([self._n_class, self._n_features + 1])
170
171     for k, class_label in enumerate(self._classes):
172         n_yk = self._class_count[k] # n samples where Y=k
173
174         X_k = self.X[self.y == class_label, :]
175
176         theta_k = n_yk / self._n_samples # P(Y=k), prior for class k
177         self._thetas[k, 0] = theta_k
178
179         for j in range(self._n_features):
180             samples_j_k = X_k[:, j] == 1 # Array with True where X_j is 1
181
182             n_xj_yk = samples_j_k.sum() # n samples where Y=k and x=x_j
183
184             if self.laplace_smoothing:
185                 theta_j_k = (n_xj_yk + 1) / (n_yk + 2)
186             else:

```

```

187         theta_j_k = n_xj_yk / n_yk
188
189         self._thetas[k, j + 1] = theta_j_k # \theta_{j, k}

```

7.3 data_processing.py

```

1 import pandas as pd
2 import numpy as np
3 import scipy.sparse as sp
4 import matplotlib.pyplot as plt
5 from sklearn.feature_extraction import text
6 from nltk import word_tokenize
7 import pickle
8 from functools import partial
9 from typing import Literal
10 import unicode
11
12 from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
13 from sklearn.decomposition import PCA
14 from sklearn.feature_selection import SelectKBest
15 from sklearn.feature_selection import f_classif, mutual_info_classif
16 from sklearn import preprocessing
17
18 import string
19 from nltk.corpus import stopwords
20 from nltk.stem import PorterStemmer
21
22 import nltk
23 from nltk.corpus import wordnet
24 from nltk import word_tokenize
25 from nltk.stem import WordNetLemmatizer
26
27 import langid
28 from langid.langid import LanguageIdentifier, model
29
30 langid.set_languages(['en', 'fr'])
31 lang_identifier = LanguageIdentifier.from_modelstring(model, norm_probs=True)
32
33 MY_STOP_WORDS = ['im', 'https', 'http', 'www', 'l', 're', 'qu', 'x200b']
34
35
36 def get_wordnet_pos(word):
37     """Map POS tag to first character lemmatize() accepts"""
38     tag = nltk.pos_tag([word])[0][1][0].upper()
39     tag_dict = {"J": wordnet.ADJ, "N": wordnet.NOUN, "V": wordnet.VERB, "R": wordnet.ADV}
40     return tag_dict.get(tag, wordnet.NOUN)
41
42
43 class LemmaTokenizer:
44     def __init__(self):
45         self.wnl = WordNetLemmatizer()
46
47     def __call__(self, doc):
48         return [self.wnl.lemmatize(t, pos=get_wordnet_pos(t)) for t in word_tokenize(doc) if t.isalpha()]
49
50
51 def MyTokenizer(text):
52     """
53     To keep $ and pound signs
54     """
55     text = text.split()
56
57     important_symbols = ['$ ', ' ', ' ', ' ']

```

```

58     for symb in important_symbols:
59         if any(symb in string for string in text):
60             text.append(symb)
61
62     return text
63
64
65 class Data:
66     def __init__(self, train_file, test_file):
67         # Download the csv data
68         self.train_file: str = train_file
69         self.test_file: str = test_file
70
71         self.train_data: pd.DataFrame
72         self.test_data: pd.DataFrame
73         self.readData()
74         self._detect_lang()
75
76         # Extract the subsets
77         self.data_list: list = self.train_data['body'].to_list() # List of all
samples
78
79         self.labels = self.train_data['label'].to_numpy() # Numpy array with label of
each sample
80
81     # Read the data
82     def readData(self):
83         self.train_data = pd.read_csv(
84             self.train_file, header=None, encoding='utf-8', skiprows=[0], names=['body
', 'label']
85         )
86         self.train_data = self.train_data.sample(frac=1, random_state=0)
87         self.test_data = pd.read_csv(self.test_file, header=None, encoding='utf-8',
skiprows=[0], names=['id', 'body'])
88
89     def _detect_lang(self):
90         """
91         To find the language (fr or en) of each post
92         """
93         self.train_data['lang'] = self.train_data['body'].apply(lambda x:
lang_identifier.classify(x)[0])
94         self.test_data['lang'] = self.test_data['body'].apply(lambda x:
lang_identifier.classify(x)[0])
95
96
97 class Format_data:
98     def __init__(
99         self,
100         words_dataset: Data, # Loaded data
101         dataset_name: str = 'NoName',
102         # Text processing options
103         max_feat: int | None = None, # Max number of tokens
104         feature_type: Literal['Bin', 'Count', 'TF'] = 'Bin',
105         n_gram: tuple = (1, 1),
106         lemmatize: bool = False,
107         lang_id: bool = False, # If true, add a feature for the language (0: en, 1:fr
)
108         rm_accents: bool = True, # To remove accents
109         standardize_data: bool = False, # To remove mean and std of all data
110         min_df: int = 1, # Ignore terms w/ frequency lower than that
111         # Feature selection options
112         feat_select: Literal['PCA', 'MI', 'F_CL'] | None = None,
113         n_feat_select: int = 1, # Number of features to keep
114         weight_samples: bool = False, # To compute the features weights
115         punc_replace: str = ' ',

```

```

116 ):
117     self.name: str = dataset_name
118     print(f"\r\tProcessing of: {self.name}... ", end='')
119
120     # Attributes
121     self.words_dataset: Data = words_dataset
122
123     # Text processing
124     self._max_feat = max_feat
125     self._n_gram = n_gram
126
127     self._feature_type = feature_type
128     if feature_type == 'Bin':
129         self._binary_features = True
130         self._use_tf_idf = False
131
132     elif feature_type == 'Count':
133         self._binary_features = False
134         self._use_tf_idf = False
135
136     elif feature_type == 'TF':
137         self._binary_features = False
138         self._use_tf_idf = True
139
140     self._lemmatize = lemmatize
141     self._lang_id = lang_id
142     self._standardize_data = standardize_data
143     self._rm_accents = rm_accents
144     self._min_df = min_df
145     self._punc_rep = punc_replace
146
147     # Feature selection
148     self._feat_select_opt = feat_select
149     self._n_feat_select = n_feat_select
150
151     # Train labels
152     self.Y = words_dataset.labels
153     self.X_test = None
154
155     self.stop_words = self._get_stop_words()
156
157     # Pre-process
158     (
159         self.train_text,
160         self.test_text,
161     ) = self._pre_process_text() # Get list of posts, lowered and w/o
162     punctuations
163
164     # Tokenize
165     self.X, self.X_test, self._vectorizer = self._vectorize_text() # _vectorizer:
166     To transform text to a vector
167
168     self.features_name = self._vectorizer.get_feature_names_out() # Corresponding
169     features of _vectorizer
170
171     self._add_lang() # Add language as a feature
172
173     self._scaler = self._normalize_data()
174
175     # Feat. Selection
176     self.pca_selector = None # PCA transformer
177     self.mi_selector = None # MI feature selection
178     self._feat_selector = self._feature_selection()
179
180     def _vectorize_text(self):

```

```

178 """
179 Create a dictionary of all words.
180
181 To get the CountVectorizer for the training dataset.
182
183 Returns:
184     _vectorizer and vectorized dataset
185
186 """
187 # Set tokenizer
188 if self._lemmatize:
189     tokenizer = LemmaTokenizer()
190
191 else:
192     tokenizer = MyTokenizer
193
194 strip_accents = 'unicode' if self._rm_accents else None
195
196 if not self._use_tf_idf:
197     vectorizer = CountVectorizer(
198         stop_words=self.stop_words,
199         max_features=self._max_feat,
200         ngram_range=self._n_gram,
201         binary=self._binary_features,
202         tokenizer=tokenizer,
203         token_pattern=None,
204         strip_accents=strip_accents,
205         min_df=self._min_df,
206     )
207
208 else:
209     vectorizer = TfidfVectorizer(
210         stop_words=self.stop_words,
211         max_features=self._max_feat,
212         ngram_range=self._n_gram,
213         binary=False,
214         tokenizer=tokenizer,
215         token_pattern=None,
216         strip_accents=strip_accents,
217         min_df=self._min_df,
218     )
219
220 # Learn the vocabulary dictionary and return document term matrix
221 X = vectorizer.fit_transform(self.train_text)
222
223 # Transform test data
224 X_test = vectorizer.transform(self.test_text)
225
226 return X, X_test, vectorizer
227
228 def _pre-process-text(self):
229     """
230     Pre-process the texts:
231     - Lowers everything
232     - Remove punctuations
233
234     Returns:
235         [str]: List with all the post preprocessed
236
237     """
238     train_df = self.words_dataset.train_data.copy(deep=True)
239     test_df = self.words_dataset.test_data.copy(deep=True)
240
241     # Lower
242     train_df['body'] = train_df['body'].str.lower()

```



```

243     test_df['body'] = test_df['body'].str.lower()
244
245     # Punctuation
246     punc_list = string.punctuation.replace('$', '')
247     punc_list += '
248
249     train_df['body'] = train_df['body'].str.replace('{}'.format(punc_list), self
._punc_rep, regex=True)
250     train_df['body'] = train_df['body'].str.replace(r'[\n\\]', '', regex=True)
251
252     test_df['body'] = test_df['body'].str.replace('{}'.format(punc_list), self
._punc_rep, regex=True)
253     test_df['body'] = test_df['body'].str.replace(r'[\n\\]', '', regex=True)
254
255     return train_df['body'].to_list(), test_df['body'].to_list()
256
257 # Specify stopwords
258 def _get_stop_words(self):
259     my_stop_words = stopwords.words('english') + stopwords.words('french')
260
261     my_stop_words += MY_STOP_WORDS
262
263     if self._rm_accents:
264         my_stop_words = [unidecode.unidecode(word) for word in my_stop_words]
265
266     # Lemmatize stop words
267     if self._lemmatize:
268         wnl = WordNetLemmatizer()
269         my_stop_words = [wnl.lemmatize(t, pos=get_wordnet_pos(t)) for t in
my_stop_words if t.isalpha()]
270         my_stop_words = list(set(my_stop_words))
271
272     # print(my_stop_words)
273     return my_stop_words
274
275 def _feature_selection(self):
276     """
277     To perform feature selection analysis
278     """
279     feat_selector = None
280     # PCA
281     if self._feat_select_opt == 'PCA':
282         pca_selector = PCA(n_components=self._n_feat_select)
283
284         if isinstance(self.X, sp.csr_matrix):
285             self.X = self.X.toarray()
286
287         self.X = pca_selector.fit_transform(self.X)
288         self.X_test = pca_selector.transform(self.X_test)
289
290         plot = True
291         if plot:
292             sing_values = pca_selector.singular_values_
293
294             # Plot the singular values
295             plt.plot(np.arange(1, len(sing_values) + 1), sing_values, marker='o')
296             plt.title(f'Singular Values - {self.name}')
297             plt.xlabel('Principal Components')
298             plt.ylabel('Singular Values')
299             plt.axvline(x=self._n_feat_select, color='red', linestyle='--', ymin
=0, ymax=1, linewidth=2)
300             plt.grid(True)
301             plt.show(block=False)
302
303         feat_selector = pca_selector

```

```

304
305     elif self._feat_select_opt == 'MI':
306         if self._use_tf_idf:
307             discrete_feat = [self.X.shape[1] - 1]
308             X = self.X.toarray()
309         else:
310             discrete_feat = True
311             X = self.X
312
313         # MI_info = mutual_info_classif(X=self.X.toarray(), Y=self.Y,
discrete_features=discrete_features, random_state=0)
314         my_score = partial(mutual_info_classif, random_state=0, discrete_features=
discrete_feat)
315         mi_selector = SelectKBest(my_score, k=self._n_feat_select)
316         self.X = mi_selector.fit_transform(X, self.Y)
317         self.X_test = mi_selector.transform(self.X_test)
318
319         selected_feats = self.features_name[mi_selector.get_support()]
320         feat_scores = mi_selector.scores_[mi_selector.get_support()]
321         names_scores = list(zip(selected_feats, feat_scores))
322         feat_scores = pd.DataFrame(data=names_scores, columns=['Feat_names', '
Score'])
323         self._feat_scores = feat_scores.sort_values(['Score', 'Feat_names'],
ascending=[False, True])
324
325         self.features_name = mi_selector.get_feature_names_out(self.features_name)
326
327         feat_selector = mi_selector
328
329     elif self._feat_select_opt == 'F_CL':
330         feat_selector = SelectKBest(f_classif, k=self._n_feat_select)
331         X_trans = feat_selector.fit_transform(self.X, self.Y)
332
333         X_test_trans = feat_selector.transform(self.X_test)
334
335         selected_feats = self.features_name[feat_selector.get_support()]
336         feat_scores = feat_selector.scores_[feat_selector.get_support()]
337         names_scores = list(zip(selected_feats, feat_scores))
338         feat_scores = pd.DataFrame(data=names_scores, columns=['Feat_names', '
Score'])
339         self._feat_scores = feat_scores.sort_values(['Score', 'Feat_names'],
ascending=[False, True])
340
341         self.features_name = feat_selector.get_feature_names_out(self.
features_name)
342
343         self.X, self.X_test = X_trans, X_test_trans
344
345     elif self._feat_select_opt is None:
346         return
347
348     else:
349         raise ValueError(f'Invalid feature selection option: {self.
_feat_select_opt}')
350
351     return feat_selector
352
353 def _add_lang(self):
354     """
355     Add a feature with the language of the post (en:0, fr:1)
356     """
357     if self._lang_id:
358         # Train
359         en_train_array = (self.words_dataset.train_data['lang'] == 'en').astype(
int).to_numpy() # 0: en,

```

```

360         en_train_array = sp.csr_matrix(en_train_array).reshape(-1, 1)
361
362         fr_train_array = (self.words_dataset.train_data['lang'] == 'fr').astype(
363         int).to_numpy() # 0: en,
364         fr_train_array = sp.csr_matrix(fr_train_array).reshape(-1, 1)
365
366         self.X = sp.csr_matrix(sp.hstack([self.X, en_train_array, fr_train_array])
367         )
368         self.features_name = np.append(self.features_name, ['is_en', 'is_fr'])
369
370         # Test
371         en_test_array = (self.words_dataset.test_data['lang'] == 'en').astype(int)
372         .to_numpy() # 0: en,
373         en_test_array = sp.csr_matrix(en_test_array).reshape(-1, 1)
374
375         fr_test_array = (self.words_dataset.test_data['lang'] == 'fr').astype(int)
376         .to_numpy() # 0: en,
377         fr_test_array = sp.csr_matrix(fr_test_array).reshape(-1, 1)
378
379         self.X_test = sp.csr_matrix(sp.hstack([self.X_test, en_test_array,
380         fr_test_array]))
381
382     def _normalize_data(self):
383         """
384         Remove mean and var of data
385         """
386
387         scaler = None
388         if self._standardize_data:
389             scaler = preprocessing.StandardScaler().fit(self.X.toarray())
390
391             self.X = scaler.transform(self.X.toarray())
392             self.X_test = scaler.transform(self.X_test.toarray())
393
394         return scaler
395
396     def get_params(self):
397         return {
398             # 'max_feat': self._max_feat,
399             'n_gram': self._n_gram,
400             'feat_type': self._feature_type,
401             'lemmatized': self._lemmatize,
402             'lang': self._lang_id,
403             'standardized': self._standardize_data,
404             'rm_accents': self._rm_accents,
405             'feat_select': self._feat_select_opt,
406             'n_feat': self._n_feat_select,
407         }

```

7.4 cross_val_score.py

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import KFold
4 import itertools
5
6 def cross_val_score(
7     model_class, X, y, cv=5, base_params=None, cv_params=None, results_df=None,
8     ds_name=None
9 ):
10     """To perform K-Fold validation to find the best combination of parameters for a
11     given model.
12
13     The parameters in 'base_params' are kept the same for all tests.

```

```

13 K-fold validation is performed for each combination of params in 'model_params'.
14
15 Args:
16     model_class (model class): Model to test
17     X (NDArray): Training dataset
18     y (NDArray): Labels of the training dataset
19     cv (int, optional): Number of folds. Defaults to 5.
20     base_params (dict, optional): Keywords to pass to the model. Defaults to {}.
21     model_params (dict, optional): Keywords to pass to the model. Defaults to {}.
22
23 Returns:
24     pd.DataFrame: Score of each combination
25 """
26 kf = KFold(n_splits=cv, shuffle=True)
27
28 results = []
29
30 # Find all combinations of parameters
31 if cv_params is not None:
32     keys, values = zip(*cv_params.items())
33     all_combs = [dict(zip(keys, v)) for v in itertools.product(*values)]
34 else:
35     all_combs = [{}]
36
37 # Find the best combination w/ CV
38 for i, each_comb in enumerate(all_combs):
39     print(f'\r\tCombination {i+1}/{len(all_combs)}', end='')
40     model = model_class(**base_params, **each_comb) # Create model with curr
41     # print(f"\tParams: {each_comb}", end='')
42
43     # Check if model has already been trained on this ds
44     if not results_df.empty:
45         matching_row = results_df[
46             (results_df['Model'].apply(type) == type(model))
47             & (results_df['Params'] == each_comb)
48             & (results_df['Dataset'] == ds_name)
49         ]
50         if not matching_row.empty:
51             continue
52
53     score = 0
54
55     comb_ok = True # Set to False if the combination of parameters is invalid
56
57     for i, (train_idx, test_idx) in enumerate(kf.split(X)):
58         if not comb_ok:
59             break
60
61         X_train = X[train_idx]
62         X_test = X[test_idx]
63         y_train = y[train_idx]
64         y_test = y[test_idx]
65         try:
66             score += model.fit(X_train, y_train).score(X_test, y_test)
67
68         except ValueError as err:
69             comb_ok = False
70             err_msg = err
71
72     score /= cv
73
74     if comb_ok:
75         # Train on whole ds
76         acc = model.fit(X, y).score(X, y)

```

```
77         results.append({'Params': each_comb, 'Score': score, 'Model': model, 'Acc'  
78         : acc})  
79  
80         if not comb_ok:  
81             print(f"Invalid model: {err_msg}")  
82  
83     return pd.DataFrame(results)
```